

CSCI-GA.3033-017 Special Topic: Multicore Programming

Homework 3

Due November 19, 2018

Please solve the following and upload your solutions to your private GitHub repository for the class as homework3.pdf by 11:59pm on the due date above. If for some reason this poses a technical problem, or you wish to include diagrams that you don't wish to spend time drawing in a drawing application, you may hand in a printed copy (*not* hand-written) at the beginning of class (6:20pm) on the day of the deadline. **Unlike labs, late homeworks will be assigned a grade of 0.**

As is university policy, instances of cheating will be taken very seriously. If you use any source for reference, including speaking with other students and/or consulting internet resources, you MUST cite those sources and/or people in your assignment. Any instances of cheating will earn you a zero on the homework, a visit to the administration, and potentially other punishments including and up to expulsion from the class and/or school.

1. Why does adding a sentinel simplify a concurrent queue? How does it make it possible to construct a lock-free queue?
2. Consider the performance of a program responsible for taking a set of 10,000 large numbers (e.g., between $1e20$ and $1e40$), and determining the prime factors of each number. A user has a four-core machine, and wants to run this program in the background while they play a fast-paced game (that spawns two threads, one for the game logic and one for rendering) at the same time.
 - a. What kind of bounds (minimum and maximum) can you give on the number of threads the prime factor program should spawn? Why not fewer or more?
 - b. What performance metric(s) would you as a programmer use? What would you optimize for in this situation, not just in your own program, but system-wide?
3. Consider the following two threads accessing shared global variable `global_thread`. (a) What kind(s) of concurrency bug(s) are present in this code? (b) **Describe** (ie, don't just list) two possible ways to fix this. If it helps, use code to demonstrate why these are good fixes.

| Thread 1 | Thread 2 |
|---|---|
| <pre>if (global_thread->proc_info) { fputs(global_thread->proc_info); }</pre> | <pre>global_thread->proc_info = nullptr;</pre> |

4.Regarding Compare And Swap (CAS):

- a. Why do we want to try to make lock-free algorithms with CAS instead of mutices or similar synchronization primitives?

b. Review our CAS-based queue `enqueue()` method, and to the best of your ability, explain why we need to allow the `tail` pointer to point to either the last or second-to-last entry in the queue instead of only the last one.

Think about the properties of CAS, what we talked about in class, and the slides. Show how you think: if you still don't quite understand it from class, you may want to explain your thought process as you try to work through examples of how two threads might interleave their `enqueue()` operations (not just the one explained in the slides). You should come up with some fundamental deductions about this and other CAS-based algorithms and the global states visible to other threads in a CAS-based algorithm that wouldn't be visible with mutex-based critical sections. Again: when in doubt, show your work!

5. What are the advantages and disadvantages of binary instrumentation versus dynamic binary translation?